CHAIR OF COMPUTER ARCHITECTURE AND PARALLEL SYSTEMS
Großpraktikum Rechnerarchitektur

Dynamic Binary Translation for RISC–V code on x86–64
Summer term 2020

Noah Dormann   Simon Kammermeier   Johannes Pfannschmidt   Florian Schmidt

Contents

1. Introduction

The aim of this project is to create an emulator capable of executing code compiled for the RISC-V instruction set architecture on an x86-64 system.

RISC-V is an open ISA first conceptualised in 2010 with the initial goals of research and education in mind. Its development took the lessons learned in terms of backwards compatibility and future-proofing from other widespread ISAs like Intel x86 into account, and aimed to provide an open interface for the architecture, rather than strict implementation details. This grants a large freedom to the implementors and greatly increases the flexibility and ease of working with the architecture [1, S. 1f].

1.1. Problem description

Because there is as yet no real hardware available for the RISC-V ISA, developers must rely on emulation in order to execute their software on a foreign architecture.

We aim to provide such an emulator, allowing the execution of RISC-V code on an x86-64 machine by means of dynamic binary translation.

By its very nature, executing code compiled for one architecture on a different one is not an easy task.

1.2. Comparison of the RISC-V and x86-64 ISAs

Continued here (compare ISAs and note challenges)...

2. Approach

In the following, the term *host* will refer to the system of the native architecture the binary translator is built for (in our case, x86-64), and the term *guest* will designate the foreign system we are attempting to emulate (RISC–V).

2.1. Modes of binary translation

When attempting to execute guest programs compiled for a foreign architecture on a different native one, there are essentially three distinct approaches at one's disposal.

The two extremes of this spectrum are:

- **Interpretation**, where, much alike interpreted programming languages (e. g. JavaScript, Python, or Ruby), the assembly instructions located in the binary are examined while emulating the execution of the program, and equivalent actions are taken on the host system in order to simulate the guest ISA.

- **Static Binary Translation**, where the guest executable is statically reverse-engineered and translated to the guest architecture as a whole. After this translation step, it can be executed as if it were a native binary, without the need for any further special treatment.

Draft after commit 10678fd9947e9f735572af684a7d74ab196f4224 on branch paper.

However, it turns out both of these approaches have their downsides. Interpretation, on the one hand, comes with a significant performance penalty. Just as programs written for interpreted languages like Python will always execute more slowly than equivalent programs created with a compiled language like C++, there is a large overhead to having to interpret every single instruction in an executable binary. Especially if the program contains certain parts of the code that are executed many times – bodies of large loop statements, for example – the interpreter has to do a lot of redundant work in simulating the same instruction operations for as many times as that section is called.

Static binary translation, on the other hand, is very difficult to achieve correctly. If done well, it would produce a binary native to the host's environment, and could thus run without any wrapper programs or other utilities, enjoying the same performance benefits as other native executables do.

However, there are certain programming concepts that have proven difficult for a static binary translator to deal with. Register indirect branching, for example, is one of them, where the jump address is not selected statically, but rather computed at runtime. Deducing the jump target would thus again require (in some ways) emulating the program execution in order to know the values needed for the computation. In light of the project at hand, a dynamic binary translator would also be a very difficult program to translate statically, due to the fact that it generates some the code it executes at runtime.

The solution to these issues, and our chosen approach for this project, is

- **Dynamic Binary Translation**, which serves as a middleground between interpreting and statically translating the executable. It aims to translate the program on the fly, while only focussing on the parts that are actually needed for execution.

By only translating the parts of the binary that are actually executed, the dynamic binary translator (DBT) can save the overhead a static translator would have spent on translating unused code paths. Also, as the DBT is actually stepping through and executing the program, the other limitations of static binary translators essentially do not apply.

It can also perform significantly better than an interpreter might, as it has to translate a single piece of code only once, which can then be called many times; unlike the above-mentioned behaviour of an interpreter. Furthermore, as it produces blocks of code that are native to the host platform, it can be reasonably expected that a single block's execution will be faster than the respective emulation.

Of course, this assumes that the translation routines of the DBT are relatively swift in performing their functions, so as not to introduce any more overhead than necessary [2, S. 1f.].

2.2. Environment setup and memory layout

As the DBT is responsible for managing the execution environment of the guest binary in the shared address space, it must also handle the setup of said environment.

The header of the ELF-file (*Executable and Linkable Format*) specifies which section(s) of the program need to be loaded, and where in memory they must reside. The DBT

must take care to map the file into memory correctly, while not compromising its own memory region.

Furthermore, the guest registers (see section 2.6 on the following page) and stack must be initialised in accordance with the architecture specification and calling convention, which necessitates a specific layout of environment and auxiliary parameters as well as command line arguments to be present [2, S. 2].

2.3. Partitioning the input code

Logically, upon facing the task of translation, the DBT must somehow divide the code into chunks it can then process for translation and execution. The natural choice here is for the translator to partition the code into basic blocks.

Basic blocks, by definition, have only a single point of entry and exit; all other instructions in a single block are executed sequentially and in the order that they appear in the code. (Of course, this does not take into account mechanisms such as out-of-order execution or system calls as well as interrupt- and exception handling).

So, for our purposes, a basic block will be terminated by any control-flow altering instruction like a jump, call or return statement, or a system call[1].

2.4. Translating the partitioned code

The most basic idea for translating the now partitioned basic blocks is to have a fixed association that maps every instruction in the guest ISA to a sequence of instructions native to the host.

The quality of the code that can be easily generated here strongly depends on the properties of the host and guest architectures in question. Difficulties can arise due to differences in the instruction operand formats and the type of instruction set architecture the DBT is dealing with.

In our case, as outlined in section 1.2 on page 3, challenges stem from the fact that we are translating code from a load-store architecture using a three-operand instruction format into a register-memory architecture in which (generally) one of the source operands is also the implicit destination operand. This, for example, means that a single arithmetic `add rd, rs1, rs2` in RISC-V assembly language generally can not be translated via a single instruction, but rather requires two instructions: moving `rs1` to `rd`, then adding the value of `rs2` to `rd`.

Opportunities for optimisation lie wherever there is a way to shorten the translation's amount of CPU clock cycles, possibly by employing semantically equivalent native instructions that run in a shorter timespan. The RISC-V pseudoinstructions (as mentioned in section 1.2) are also of some help here [1, S. 139], as it is clear that an instruction like `xori x10, x10, -1` can be directly translated as a `not x10`, without needing to resort to `mov` and `xor`.

---

[1]These may or may not have control-flow altering effects; they in any case need to be handled this way due to the reasons laid out in section 2.7 on page 7.

2.5. Code cache and block handling

Naturally, the DBT aims to store the translated code in a semi-permanent way, for it is the goal to not have to translate a required section more than once.

For that, we allocate a region of memory reserved for the basic block translations, also called a *code cache*. Additionally, an index to this memory section is required, since there needs to be a way to quickly reference the blocks residing in the cache and associate them with both the host and guest instruction pointers that identify them during execution.

It is possible that this code cache might fill up during the execution of a large guest program. If it does, there are two different strategies to handle this issue: One can either invalidate and purge some or all of the blocks currently residing in the cache, or dynamically resize the cache according to the needs of the guest program [2, S. 3].

Purging the entire cache would require the translator to restart translation on older blocks that might be needed again, introducing a performance overhead that needs to be weighed against the higher memory usage of enlarging the cache.

On the other hand, selective deletion of some of the blocks in the cache is very difficult due to optimisations taken in the context of chaining. As any chained jumps located in another cached block are dependent on the target block residing in the cache, the target's removal would invalidate these jumps. It would thus only be possible to either remove all blocks with jump references to the candidate up for removal, or to leave all blocks with jump references in the cache altogether.

2.6. Register handling and context switching

As outlined in section 1.2, the RISC-V and x86-64 architectures have differing amounts of general purpose registers. In some way, the state of the 32 general purpose registers $x1$[2] to `x31` and the `pc` needs to be stored and available to the translations of the identified basic blocks.

As x86-64 only provides for 16 general-purpose registers (`rax-rdx, rsi, rdi, rsp, rbp and r8-r15`), it is impossible to directly and statically map all guest registers to native host registers. Adding to the above, due to the fact that some x86-64 registers have special or implicit purposes in some instructions (`rax` and `rdx` in `(i)mul`, `cl` in shifting, etc.), care must be taken in choosing the registers that can be used for such a mapping. Keeping a guest register file exclusively in memory, and loading them into native registers when needed within the translations of single instructions is technically possible, especially in light of x86-64's ability to extensively use memory operands in the instructions. However, this necessitates a large number of memory accesses for both memory operands in the instructions as well as local register allocation within the translated blocks. Due to the very large performance gain connected to using register operands instead of memory operands, this is also not feasible at scale [2, S. 8f.].

Accordingly, the solution would be an approach that employs parts of both of these extremes [2, S. 9]. We aim to employ the tools we design to discover the most-used

---

[2] `x0` is hardwired to a constant zero. All reads will return 0, all writes will be ignored. Hence, this register needs special handling in the DBT.

registers in the guest programs, and statically map these to general purpose x86-64 regis-ters. The remaining operands can either be used from memory directly, or dynamically allocated into free host registers inside a single block's translation. Thus, we save much overhead otherwise spent on memory access to the register file, but do not unnecessarily occupy native register space with seldomly accessed guest registers.

2.7. System call handling

System calls are also a very important part of enabling the guest program's execution. Thus, every ISA must offer some way to switch the execution context in the kernel mode for the system call to be handled.

For RISC-V, the instruction `ECALL` (for *environment call*, formerly `SCALL`) handles these requests, with the system call number residing in register `a7` and the arguments being passed in `a0-a6`.

However, the DBT generally cannot just reorganise the guest argument values and system call identifier according the host's calling convention and relay the system call directly. The RISC-V guest program expects a different operating system kernel than is present natively on the host; with that, the system call interface also differs [2, S. 2f.].

In order to handle the `ECALL` instruction correctly, the translator must thus build the translated instruction to call a specific handler routine not too dissimilar from one that may be found in a kernel. There, system calls that exist natively on the host architecture as well (like `write` or `clock_gettime`) can usually be passed along to the host's kernel directly.

Care must be taken for system calls that would enable the guest to change the state or context of the host – an `mmap` into the translator's memory region, for example, or a call to `exit` – these calls must be emulated accordingly to prevent these faults. In cases where the data structure layout used by the kernels differs, the DBT must also perform necessary actions to adapt the formats to each other. Some system calls may not exist at all on the native architecture of the host, it is up to the DBT to emulate the required functionality [2, S. 2f.].

3. Implementation Details

The following section aims to provide an in-depth overview of the system's architecture as well as the rationale for major design decisions taken during the implementation.

3.1. System architecture and execution control flow

3.2. Instruction translation process

3.3. Code cache and TLB for block lookup

3.4. Static hybrid register mapping

3.4.1. Register priority analysis

In order to achieve the best performance with the hybrid approach to the register mapping described in section 2.6 on page 6, we must decide which registers of the RISC-V guest to map into the host's limited number of available GPRs. There are two main ways of determining the priority of registers when considering them as candidates for a mapping.

It is, on the one hand, possible to assess the priority statically, by performing an analysis of the binary in question. Essentially, the hereby produced metric counts the number of times the register is used in the assembly instructions listed in the guest program and thus delivers an idea of how important each register is to this specific executable. We have built the tools required for this effort directly into the translator's analyser function, accessed via the -a flag (see appendix D).

However, this approach does not take into account that a single instruction may be executed many times while the program is running. Accordingly, the other approach is to assess the register priority dynamically by analysing and profiling the execution of the testing program, thereby gaining an insight into how often each register is actually used during the execution. The translator is also capable of performing such an analysis, commanded by the -p flag. A dynamic analysis, of course, delivers a largely more accurate idea of the priority of the registers in question, but has the decided and obvious disadvantage that it cannot be performed without actually executing the binary.

For the results of such an analysis performed on a range of programs, including *gzip* [3] and several benchmarks of the intspeed-Suite of *SPEC CPU 2017* [4], see table 1 on the following page. Primarily, we gain interesting insights into the differences between the static and dynamic results yielded by the analysis. While the static ranked hit list does not differ greatly between the different executables and the top 12 entries are identical for every one of the tested programs, the dynamic results are far more variable. This makes creating a register mapping that fits well to every executable very difficult.

The benchmarks 605.mcf_s (route planning workload) and 620.omnetpp_s (discrete event simulation for computer networking) [5] of the *SPEC CPU* suite can serve as examples here. For programs like 605.mcf_s that only lightly use the stack, holding the stack pointer sp/x2 in a native register when only $1{,}20\,\%$ of accesses actually utilise it would not be necessary. However, other programs like 620.omnetpp_s may rely heavily on the stack, and thus log very frequent accesses to sp/x2; when statistically every ninth access is to the stack pointer, it is absolutely essential to map the register to a native GPR.

If a static analysis yielded results of similar quality to the dynamic counterpart, the DBT could analyse the binary prior to execution and run every program with a best-fit static register mapping. However, evidently, this is impossible with dynamic profiling.

Register analysis results

*: for SPEC benchmarks: test workload was chosen
**: workload: approx. 150 MB text file sourced from /dev/urandom

Static: Number of times the register is mentioned in the binary. Computed statically by parsing and analysing the **Run via:** ./translator -a -f <binary>
Dynamic: Number of register accesses during execution. Computed dynamically by logging each read/write to a pro **Run via:** ./translator -p -f <binary> [guest arguments]

Table 1: The results of a static and dynamic register usage analysis of several *SPEC CPU 2017* benchmarks, *gzip* as well as a merge sort utility program.

Static analysis — Average Static Hits: 77,32% mapped, 22,68% not mapped

Dynamic — Average Dynamic Hits: 83,59% mapped, 16,41% not mapped

RISC-V register	a5	a4	a3	a0	fp	sp	a2	a1	s1	ra	a7	s2
x86-64 mapping	rbx	rbp	rsi	rdi	r8	r9	r10	r11	r12	r13	r14	r15

Table 2: The static register mapping in use by the translator.

3.4.2. Structure of the mapping

When we structure our mapping by the average case of the insights gained, we statistically capture about $83{,}59\,\%$ of register accesses, leaving the remaining $16{,}41\,\%$ to read from the register file in memory.

From the 16 general-purpose registers x86-64 has to offer, we may use the 12 registers rbx, rbp, rsi, rdi and r8–r15. The remaining registers have either implicit or exclusive functions in some instructions (rax and rdx for multiplication/division, cl for shifting), or, like rsp, are impractical to use in combination with block chaining and function calls.

Taking the 12 registers that are most accessed on average, the mapping structure is as seen in table 2.

3.4.3. Register file memory access

Continued here. . .

3.5. Context switching details

3.6. Optimisation of the generated code

3.6.1. Block chaining

3.6.2. Recursive jump translation

3.6.3. Return address stack

3.6.4. Macro operation fusion by pattern matching

3.7. Detailed system call overview

As described in section 2.7 on page 7, we must assume the role of the kernel by handling system calls during the execution of the guest program. We achieve this by translating the ECALL instruction as a context switch and jump to the emulate_ecall routine in the DBT, which can then take the appropriate action.

As we stored the guest's registers before jumping to the handler, the requested system call index is now available to the DBT in the register file as entry a7, as per the RISC-V standard calling convention. We may now handle the system calls based on that index and the arguments passed in the registers a0 through a6, and write the return value to entry a0 of the register file prior to switching the context back to the guest.

As previously mentioned, some system calls require special handling when encountered by the DBT (see table 3 on page 12 for details). The following will describe the

specifics of these issues with system calls that are either not present on the x86-64 host architecture, or may influence or break the state of the DBT.

Adapting structure data format. There are system calls like `fstat` and `fstatat` that exist both on RISC-V as well as x86-64, but use different data structure layouts in their return values. Thus, the DBT must adapt the host's returned data to the required format prior to passing it back to the guest.

Emulation required. The DBT captures the `exit` and `exit_group` calls. Passing them through would immediately terminate the DBT – an action that is undesirable as it prevents any form of clean-up or post-execution profiling and analysis to take place. Thus, the DBT uses these system calls to set a flag which stops the translator's main loop from executing the next iteration.

The `brk` system call must also be entirely emulated, as it would otherwise allow the guest program to modify the endpoint of the DBT's data segment (*program break*), thus potentially deallocating some of the translator's memory.

Ignoring system calls. The `rt_sigaction` system call is ignored by the DBT, as any signals received by the process will be handled by the translator due to the fact that the guest program's execution is emulated in the DBT's process.

Guarded pass-through to host. Essentially, any system call that has the possibility to influence the state or memory of the translator needs to have respective safe-guards in place. A good example of this behaviour is the `mmap` system call, the handling of which also reflects the memory layout scheme discussed in section 2.2 on page 4.

In any case, we must prevent a memory mapping into the translator's memory region. Mappings that do not interfere with the DBT's memory can be passed along to the host directly. In case a hinted mapping would conflict with the translator's memory, we may just re-hint the mapping to the top of the guest's address space. When the call is not hinted (the `MAP_FIXED` or `MAP_FIXED_NOREPLACE` flag commands the mapping at exactly the specified address), we are unable to provide the guest with the requested mapping; thus we simulate an existing mapping in the location in question by returning `EEXIST` for `MAP_FIXED_NOREPLACE` and failing the call with `EINVAL` for `MAP_FIXED`.

Similarly, we fail the guest's `munmap` with `EINVAL` in cases where the translator's memory would be compromised by the deallocation.

The other supported system calls may be directly passed through to the host after performing the necessary index mapping (see table 4 on page 13).

4. Results and Performance

Measuring the performance of the DBT was accomplished by using the tools in *SPEC CPU 2017*'s `intspeed` suite of benchmarks. This not only generates reproducible and

System Call (index)	Handling	x86-64 base (index)
`fstatat` (79)	data reformat	`newfstatat` (262)
`fstat` (80)	data reformat	`fstat` (5)
`exit` (93)	emulate	n/a
`exit_group` (94)	emulate	n/a
`rt_sigaction` (134)	ignore	n/a
`brk` (214)	emulate	n/a
`munmap` (215)	guarded pass-through	`munmap` (11)
`mmap` (222)	guarded pass-through	`mmap` (9)

Table 3: An overview of the system calls we support that require special handling by the binary translator.

widely accepted results in the industry, it also validates the results produced during the run, thus ruling out any errors in the benchmark's translation.

The `intspeed` suite also presents a variety of different workloads to the translator that are based on real-life scenarios, thus producing an accurate and understandable overview of the DBT's performance in a non-controlled environment. Further context is provided by performance testing using the data compression utility *gzip* [3], where compression time is compared between runs on a native machine, in QEMU and in the DBT.

All testing was performed on an x86–64 8-core *Intel Xeon Bronze 3106* system clocked at 1,70 GHz base with 78 GiB of physical memory, running *Ubuntu 18.04.3 LTS*, kernel version *4.15.0-70-generic*. The DBT was compiled via `CMAKE_BUILD_TYPE` set to `Release`, which implies `-O3`.

4.1. SPEC CPU 2017 Results

4.2. Data compression via gzip

Next to the results of the *SPEC CPU 2017* suite, it is also valuable to measure the performance of the translator in real-world workloads by running data compression via *gzip*.

The native *gzip* binary for the reference runs was obtained from the default package repositories. As there is no such executable available for the RISC-V architecture, the sources had to be manually compiled.

Figure 2 on page 15 lists the execution times of *gzip* compressing a pseudo-random 500 MB file sourced from `/dev/urandom`[3].

---

[3]Reproducible via `base64 /dev/urandom | head -c 524288000 > random.txt;`

RISC-V system call	. . . index	x86-64 index
getcwd	17	→ 79
fcntl	25	→ 72
ioctl	29	→ 16
unlinkat	35	→ 263
faccessat	48	→ 269
chdir	49	→ 80
fchmod	52	→ 91
fchown	55	→ 93
pipe2	59	→ 293
openat	56	→ 257
close	57	→ 3
getdents64	61	→ 217
lseek	62	→ 8
read	63	→ 0
write	64	→ 1
writev	66	→ 20
readlinkat	78	→ 267
utimensat	88	→ 280
set_tid_address	96	→ 218
futex	98	→ 202
set_robust_list	99	→ 273
clock_gettime	113	→ 228
tgkill	131	→ 234
rt_sigprocmask	135	→ 14
uname	160	→ 63
gettimeofday	169	→ 96
getpid	172	→ 39
getuid	174	→ 102
geteuid	175	→ 107
getgid	176	→ 104
getegid	177	→ 108
gettid	178	→ 186
sysinfo	179	→ 99
execve	221	→ 59
wait4	260	→ 61
prlimit64	261	→ 302
renameat2	276	→ 316
getrandom	278	→ 318

Table 4: An overview of the system calls handled via pass-through to the host.

SPEC CPU 2017

Figure 1: Results of `ref`-workload runs of *SPEC CPU 2017*'s `intspeed`.

4.3. Other performance metrics

By implementing optimisations to the DBT like a return address stack, recursive jump target translation and block chaining as discussed in section 3.6 on page 10, we are able to meaningfully increase the performance of the translator in certain workloads, to the point where we outperform *QEMU* significantly.

Figure 3 on page 16 shows a manifestation of these performance gains, as seen by executing a benchmark based on a recursive implementation of the merge sort algorithm.

5. Summary

Summary here. . .

Random Data Compression (*gzip*, 500 MB, 10 runs)

Figure 2: Execution time of gzip file compression (500 MB of random data, 10 runs) in seconds (lower is better).

Appendices

A. Download and installation instructions

The source code for the translator can be downloaded by checking out the project's git repository. Take care to either `git clone` the repository with the option flag `-recursive`, or to run the command

```
1   git submodule update --init
```

as the repository contains submodules that are required for compilation.
    Then, the translator can be built by exeucting

```
1   sudo apt-get -y install gcc g++ cmake make autoconf meson
2   mkdir build && cd build && cmake .. && make
```

in the root directory of the repository. Note that the build requires CMake version 3.16 or above. This will build two artifacts:

translator The actual dynamic binary translator. For details on the usage, see section C on the next page, or execute `./translator -h`.

test The unit test binary. It can be executed via `./test` and performs extensive unit testing of the RISC-V instruction implementations, the cache, register file, as well as the parser.

Figure 3: Execution time of merge sort (5000 runs) in nanoseconds (lower is better). Testing performed with v1.2.2-61.

B. Executable program requirements

We can execute binaries compiled via the tools provided in the RISC-V GNU toolchain[4]. The executables need to be linked statically (pass the flag `-static` to gcc when compiling), as the translator does not support dynamically linked files.

We currently support binaries compiled for the architecture specifier `rv64ima`, meaning the compiler is free to utilise the base integer instruction set (`i`), as well as the instructions provided by the multiplication (`m`) and atomic standard extensions (`a`). This can be achieved by passing `-march=rv64ima` to gcc[5].

C. Using the translator

```
1    ./translator [translator option(s)] -f <filename> [guest options]
```

Seen above is the syntax for executing the translator with a guest program. All possible translator options are described in the help text, as seen by executing `./translator -h`. Every option after the filename specified via the `-f` flag is passed along to the guest in its `argv`, so all options intended for the translator must be passed before `-f`.

The command line options also include the ability to analyse (`-a`) the binary to produce a detailed breakdown of which instruction mnemonics and registers the guest will use when executed. Furthermore, it includes the ability to time the execution of only the guest program by passing the flag `-b`.

---

[4]For further information as well as download and usage instructions, see `https://github.com/riscv/riscv-gnu-toolchain` (last accessed on 25.09.2020).

[5]Note that some architecture strings require recompilation of the toolchain. Also, the aforementioned option implies `-mabi=lp64`.

Logging can be controlled by passing the requested category to the `--log` flag as detailed in the help, and can provide insights into the state of the translator during execution or debugging. Lastly, it is also possible to selectively disarm optimisation features like the return address stack, block chaining or recursive jump target translation via `--optimize`.

D. Using the bundled tools

There seems to be nothing here yet...

E. Version history

The following mirrors the version control change log of the translator over the course of its development.

Version 1.2.2 (latest)

- expand the static register mapping for better performance overall
- reallocate the code cache index and rehash all values when it is 50 % full for better lookup performance on capacity overflow
- rewrite command line options parsing to allow for long options (see `./translator -h`)
- allow finer control of specific optimisation features via `--optimize`
- add instruction pattern matching to the binary analyser to gather data for macro operation fusing
- add a profiler for counting register accesses
- implement emulation for syscalls `faccessat, getrandom, renameat2`
- remove emulation for the syscall `clone`
- fix crash when the code cache fills up by increasing the memory space available for translated blocks

Version 1.2.1

- add implementations for `AMOMIN` and `AMOMAX` instructions
- add extensive unit testing for atomic and arithmetic instructions, as well as the parser
- fix several issues to enable the SPEC CPU 2017 benchmark suite to run
- implement emulation for syscalls `chdir, pipe2, getdents64, munmap, clone, execve, wait4`
- fix `ORI` instruction being parsed as `XORI`
- fix instruction semantics for `SUB(W)`
- finalize implementation of the return address stack

Version 1.2.0

- enable register mapping for translated instructions
- add context switching from host to guest programs
- rework instruction translator function for flexibility
- implement a return address stack
- implement a TLB for cache lookup of blocks
- flip `-m` translation optimiser flag (enabled by default, flag now turns off optimisations)

Version 1.1.0

- cleanup and refactor project files
- remove all C++ usage from translator code
- eliminate standard library usage
- add performance measuring flags to the translator

Version 1.0.1

- fix an issue with the `read` system call that causes blocking problems with gzip

Version 1.0   The initial release of the translator capable of executing gzip. This release supports

- the RISC-V integer instruction set
- the multiplication extension instructions (M)
- the atomic extension instructions (A).

The latter are not yet implemented atomically, however they make the translator compatible with binaries compiled for the architecture `rv64ima`, with the ABI `lp64`.

List of Tables

List of Figures

References

[1] Editors Andrew Waterman and Krste Asanović, *The RISC-V Instruction Set Manual, Volume I: User Level ISA, Document Version 20191213*. RISC-V Foundation, Dec. 2019.

[2] M. Probst, "Dynamic binary translation," in *UKUUG Linux Developer's Conference*, vol. 2002, 2002.

[3] "The gzip home page." `https://www.gzip.org/` (last visited 02.10.2020), 2020.

[4] "SPEC CPU 2017." `https://www.spec.org/cpu2017/` (last visited 02.10.2020), 2020.

[5] "SPEC CPU 2017 Documentation." `https://www.spec.org/cpu2017/Docs/overview.html` (last visited 02.10.2020), 2020.